

Towards Developing High Performance RISC-V Processors

Using Agile Methodology

Yinan Xu*†, Zihao Yu*, Dan Tang*‡, Guokai Chen*†, Lu Chen*†, Lingrui Gou*†, Yue Jin*†, Qianruo Li*†, Xin Li*†, Zuojun Li*†, Jiawei Lin*†, Tong Liu*, Zhigang Liu*, Jiazhan Tan*, Huaqiang Wang*†, Huizhe Wang*†, Kaifan Wang*†, Chuanqi Zhang*†, Fawang Zhang II, Linjuan Zhang*†, Zifei Zhang*†, Yangyang Zhao*, Yaoyang Zhou*†, Yike Zhou*, Jiangrui Zou II, Ye Cai II, Dandan Huan¶, Zusong Li¶, Jiye Zhao¶, Zihao Chen§, Wei He§, Qiyuan Quan§, Xingwu Liu, Sa Wang*†, Kan Shi*, Ninghui Sun*† and Yungang Bao*†**

**State Key Lab of Processors, Institute of Computing Technology, Chinese Academy of Sciences*

†University of Chinese Academy of Sciences

‡Beijing Institute of Open Source Chip

§Peng Cheng Laboratory

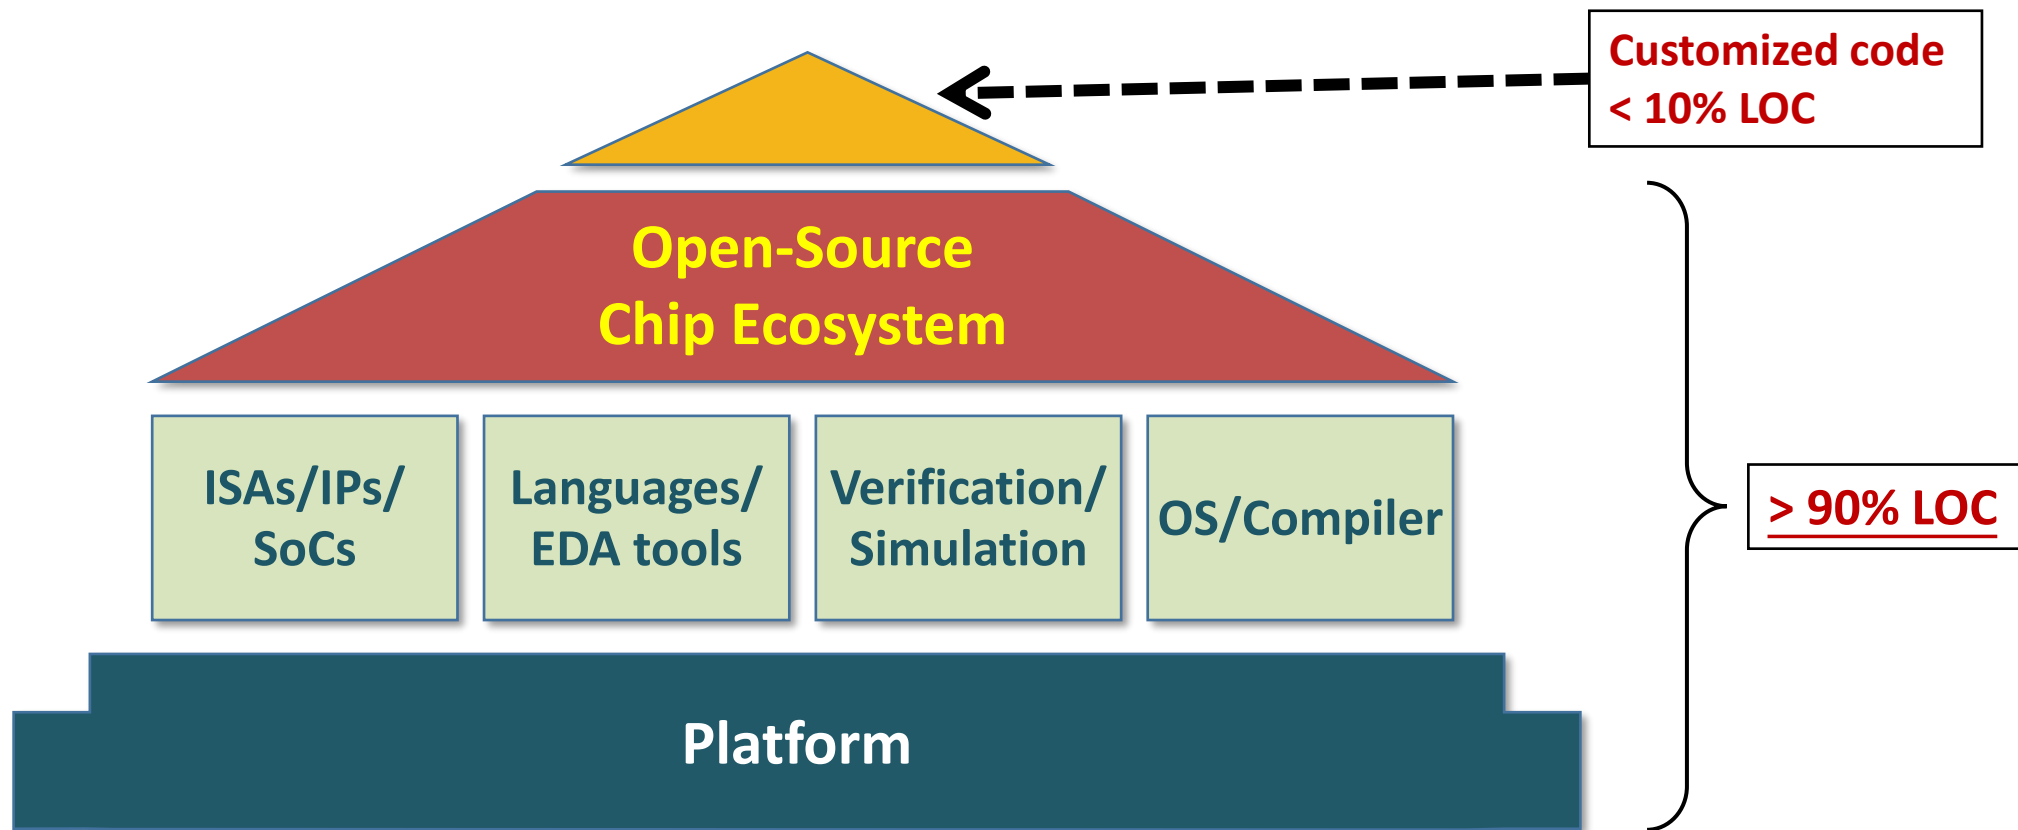
¶Beijing VCore Technology

//Shenzhen University

***Dalian University of Technology*



The Era of Agile and Open-Source Hardware



To lower the barrier of chip development
by saving time-to-market and the cost of
IPs, EDA tools, facilities and engineers etc.

🏔 Agile and Open-Source Chip Ecosystem

Customized code < 10% LOC

Open-Source Chip Ecosystem

Open-Source
> 90%

 **RISC-V**



XIANGSHAN

Rocket Chip Generator 🚀

OpenPiton



BlackParrot


ISAs/IPs/SoCs

CHISEL

bluespec



SpinalHDL



PyMTL

OpenROAD

Languages/EDA tools

 **imperas**



VERILATOR

ESSENT:

A High-Performance RTL Simulator

 **FireSim**

Verification/Simulation

Linux™ 

 **RT-Thread**



GCC



OS/Compiler

Platform

Agile Approaches Adopted in the Industry

- When we talked to some of the leading companies ...

We:



ISAs/IPs/
SoCs

Languages/
EDA tools

Verification/
Simulation

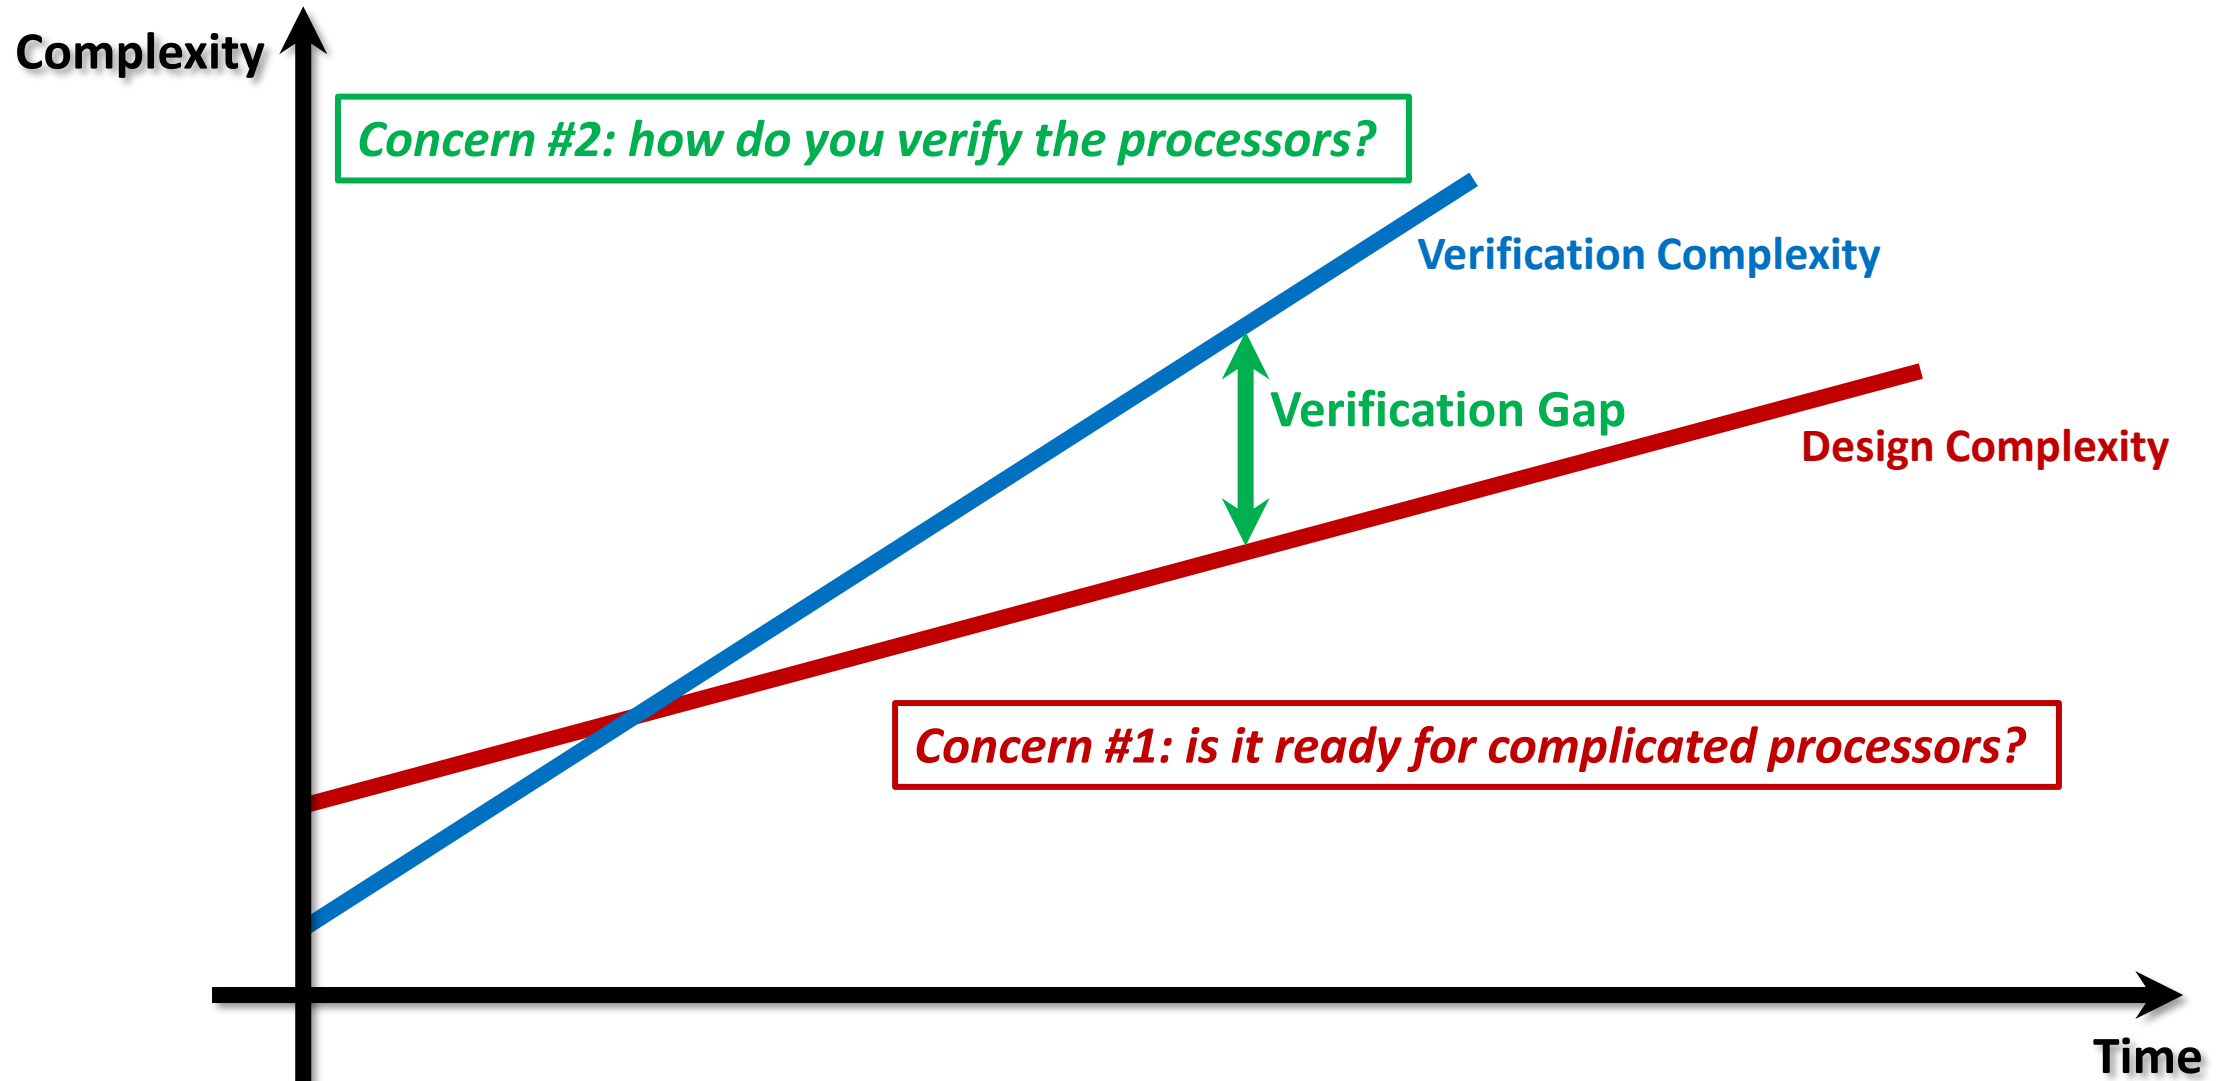
OS/Compiler

Big Ones:





Major Concerns Regarding the Agile Methodology





This Work: Let's *Do It* and See What's Happening

We:



ISAs/IPs/
SoCs

Languages/
EDA tools

Verification/
Simulation

OS/Compiler


Big Ones:



Concern #1: it's not ready for complicated processors.

Concern #2: the verification process is still less agile.

We:



XiangShan: High Performance RISC-V Processors

LET'S DO THIS THING!



XiangShan: Open-Source High Performance Processors


• 1st generation: YQH

- RV64GC, single-core, superscalar OoO
- 28nm tape-out, 1.3GHz, July 2021
- SPEC CPU2006 7.01@1GHz, DDR4-1600



• 2nd generation: NH

- RV64GCBK, dual-core, superscalar OoO
- Scheduled 2GHz@14nm tape-out, Q4 2022
- Estimated** SPEC CPU2006 19.45@2GHz



• 3rd generation: KMH

- RV64GCBKHV, quad-core, superscalar OoO
- Close collaboration with industrial partners



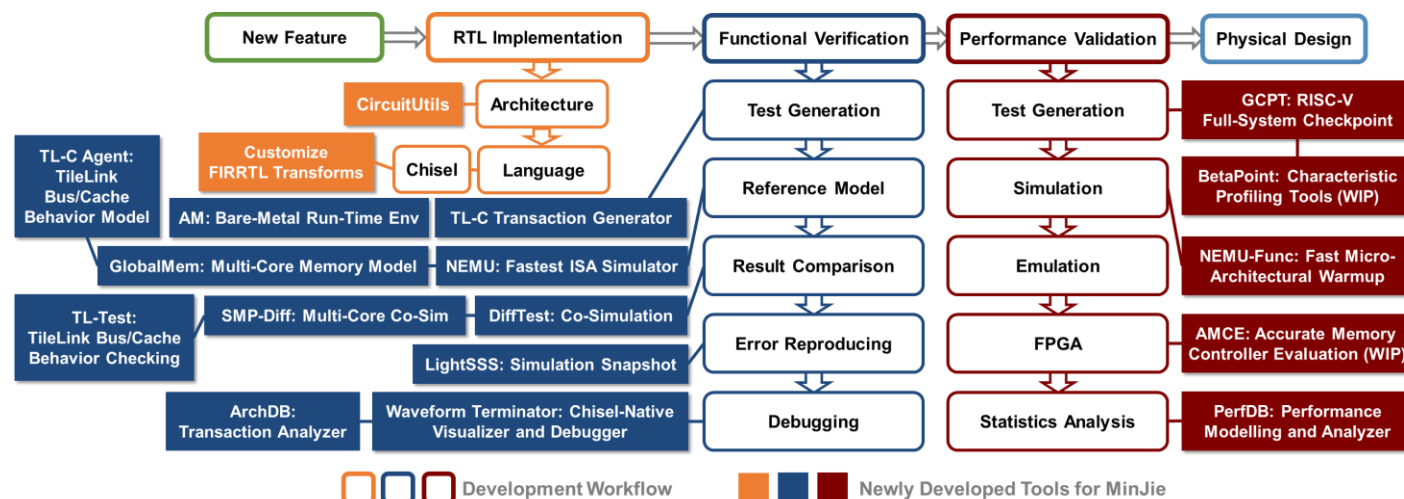




* Source: XT910@ISCA'20, SiFive, AnandTech

** Updated September 14, 2022

How We Built XiangShan: MinJie



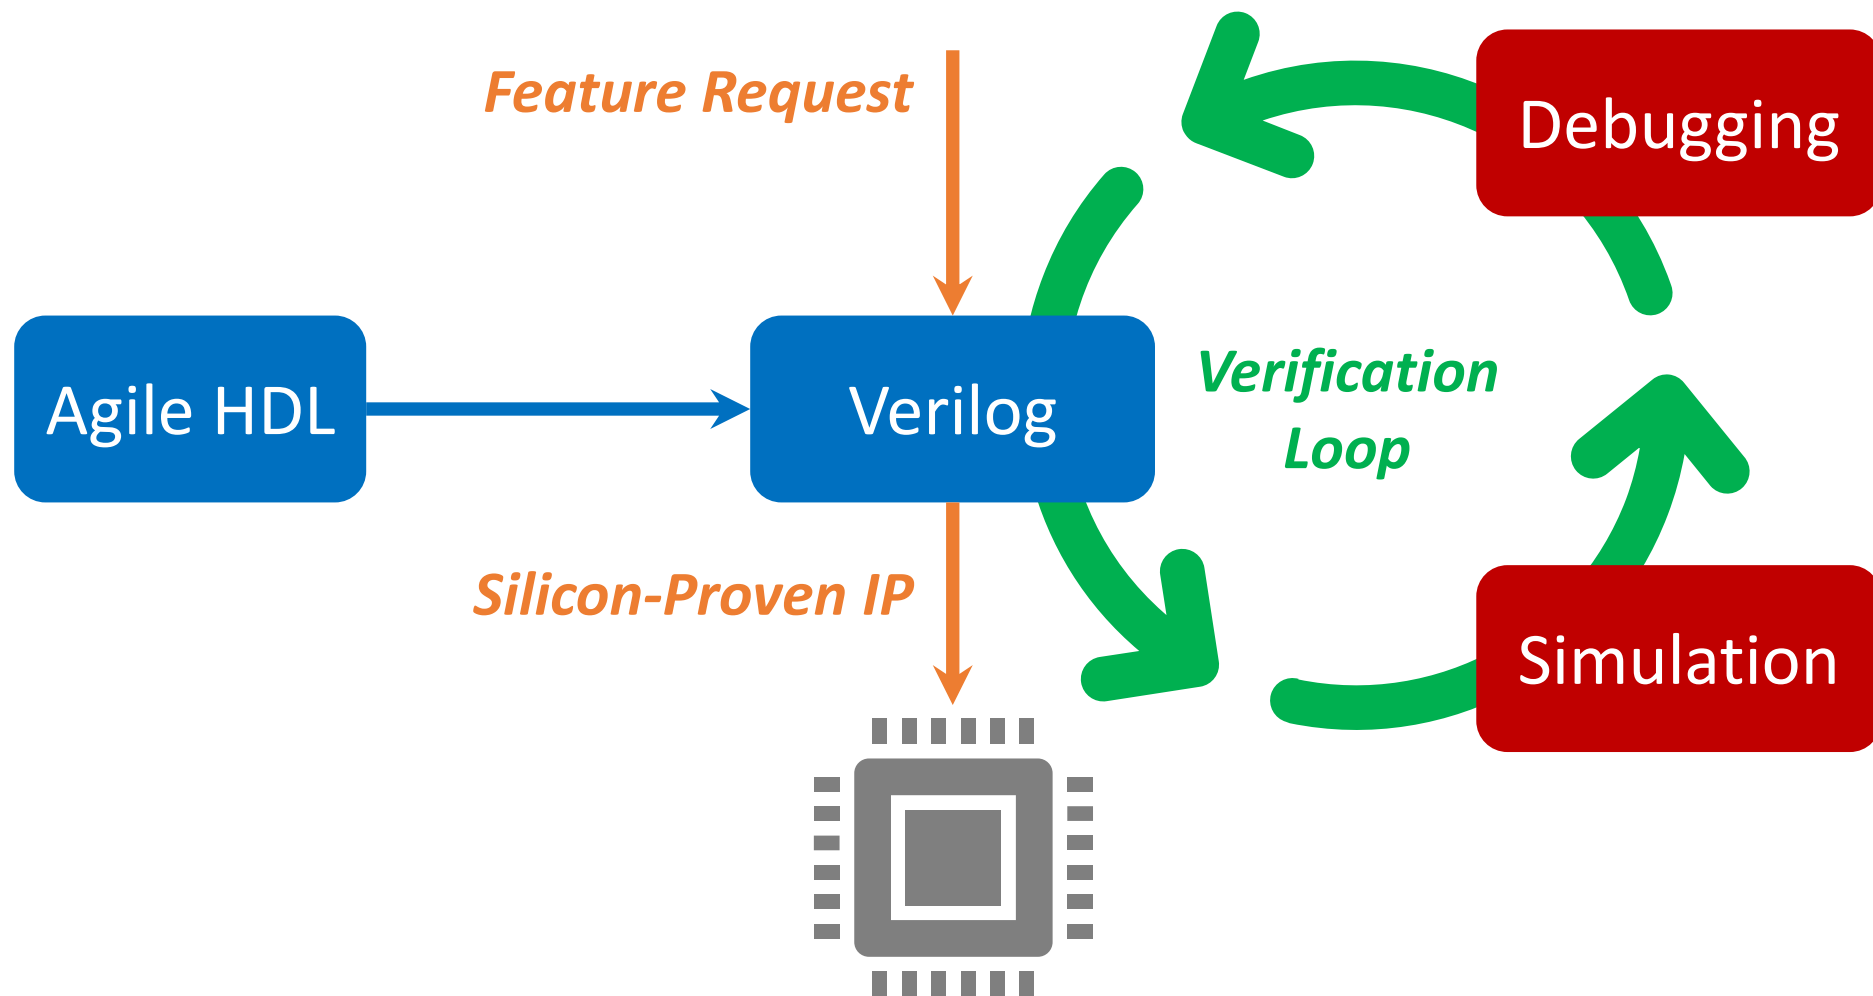
XiangShan: Open-Source High Performance RISC-V Processor



MinJie: Open-Source Platform with Agile Development Flows and Tools

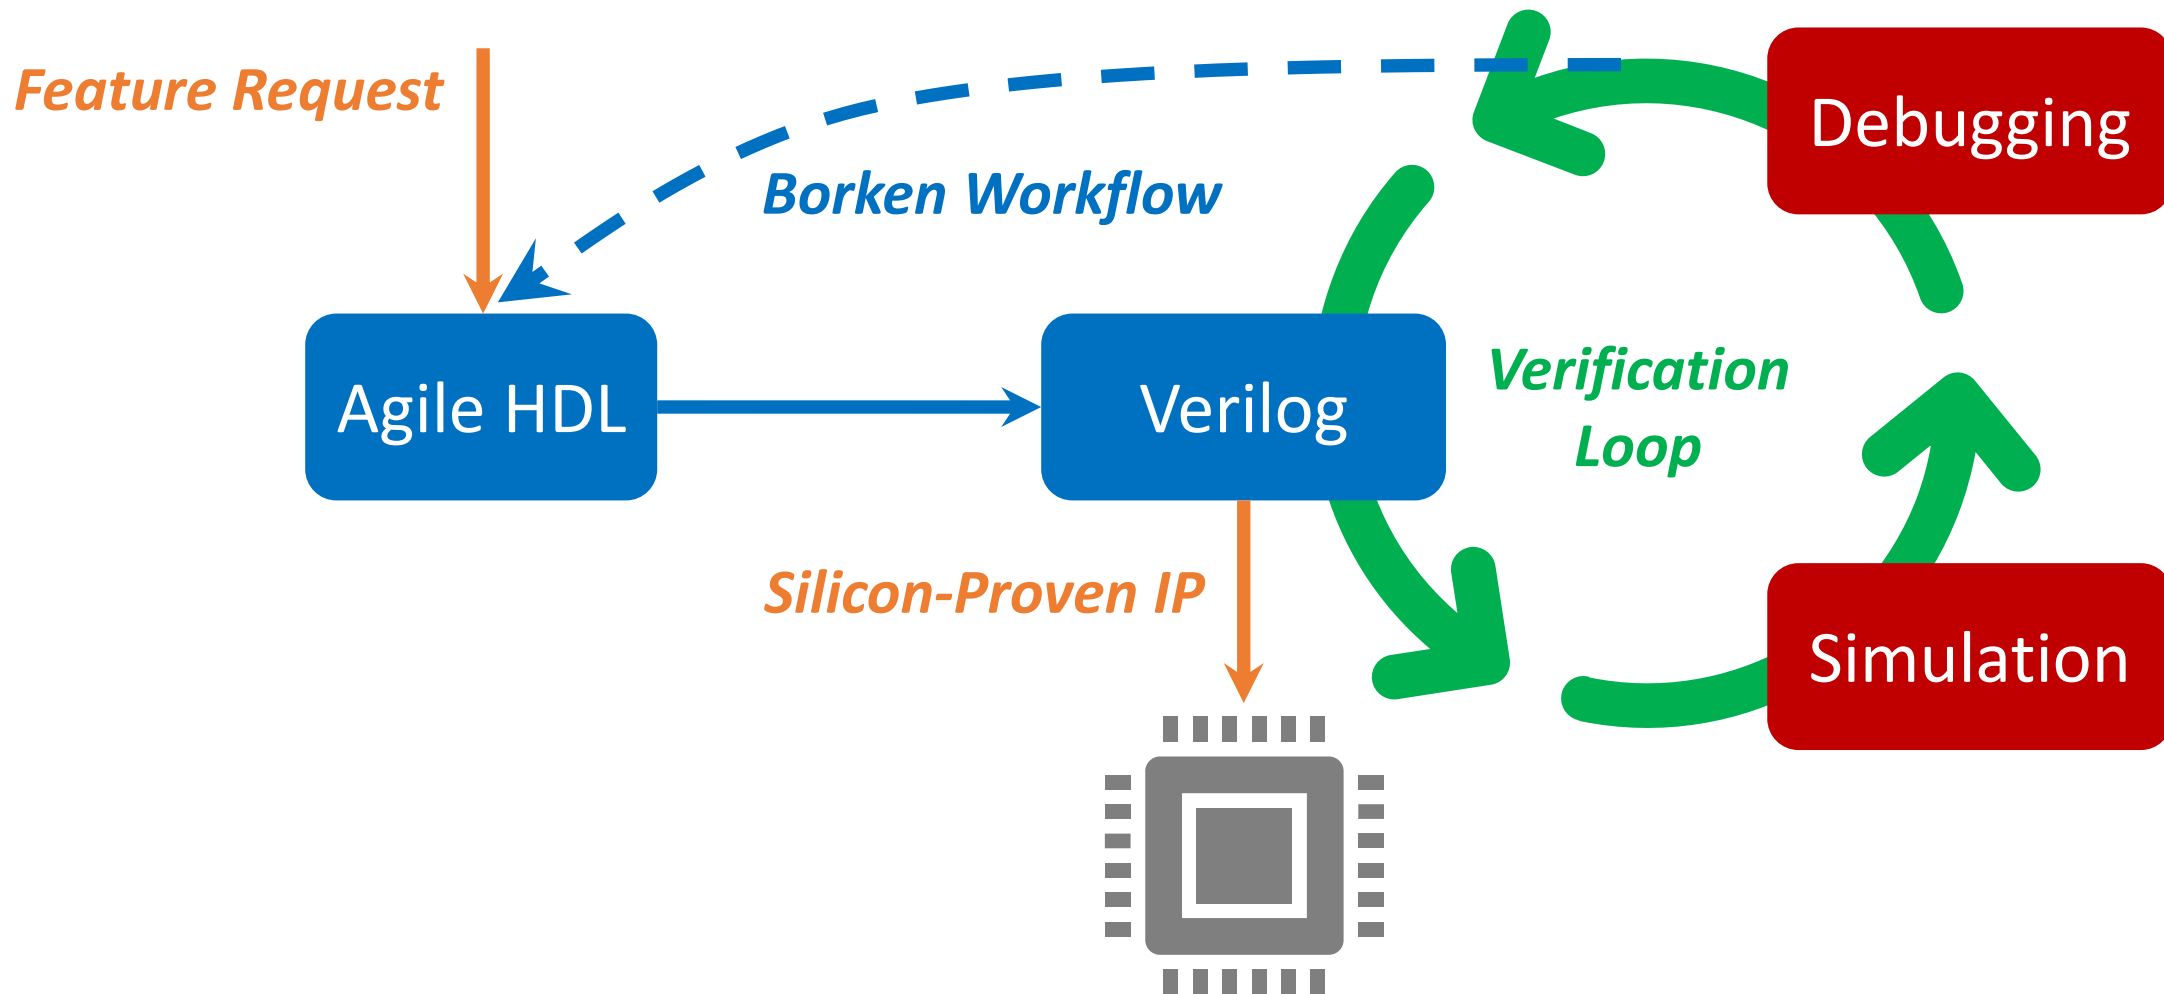


MinJie: Platform with Agile Development Flows and Tools



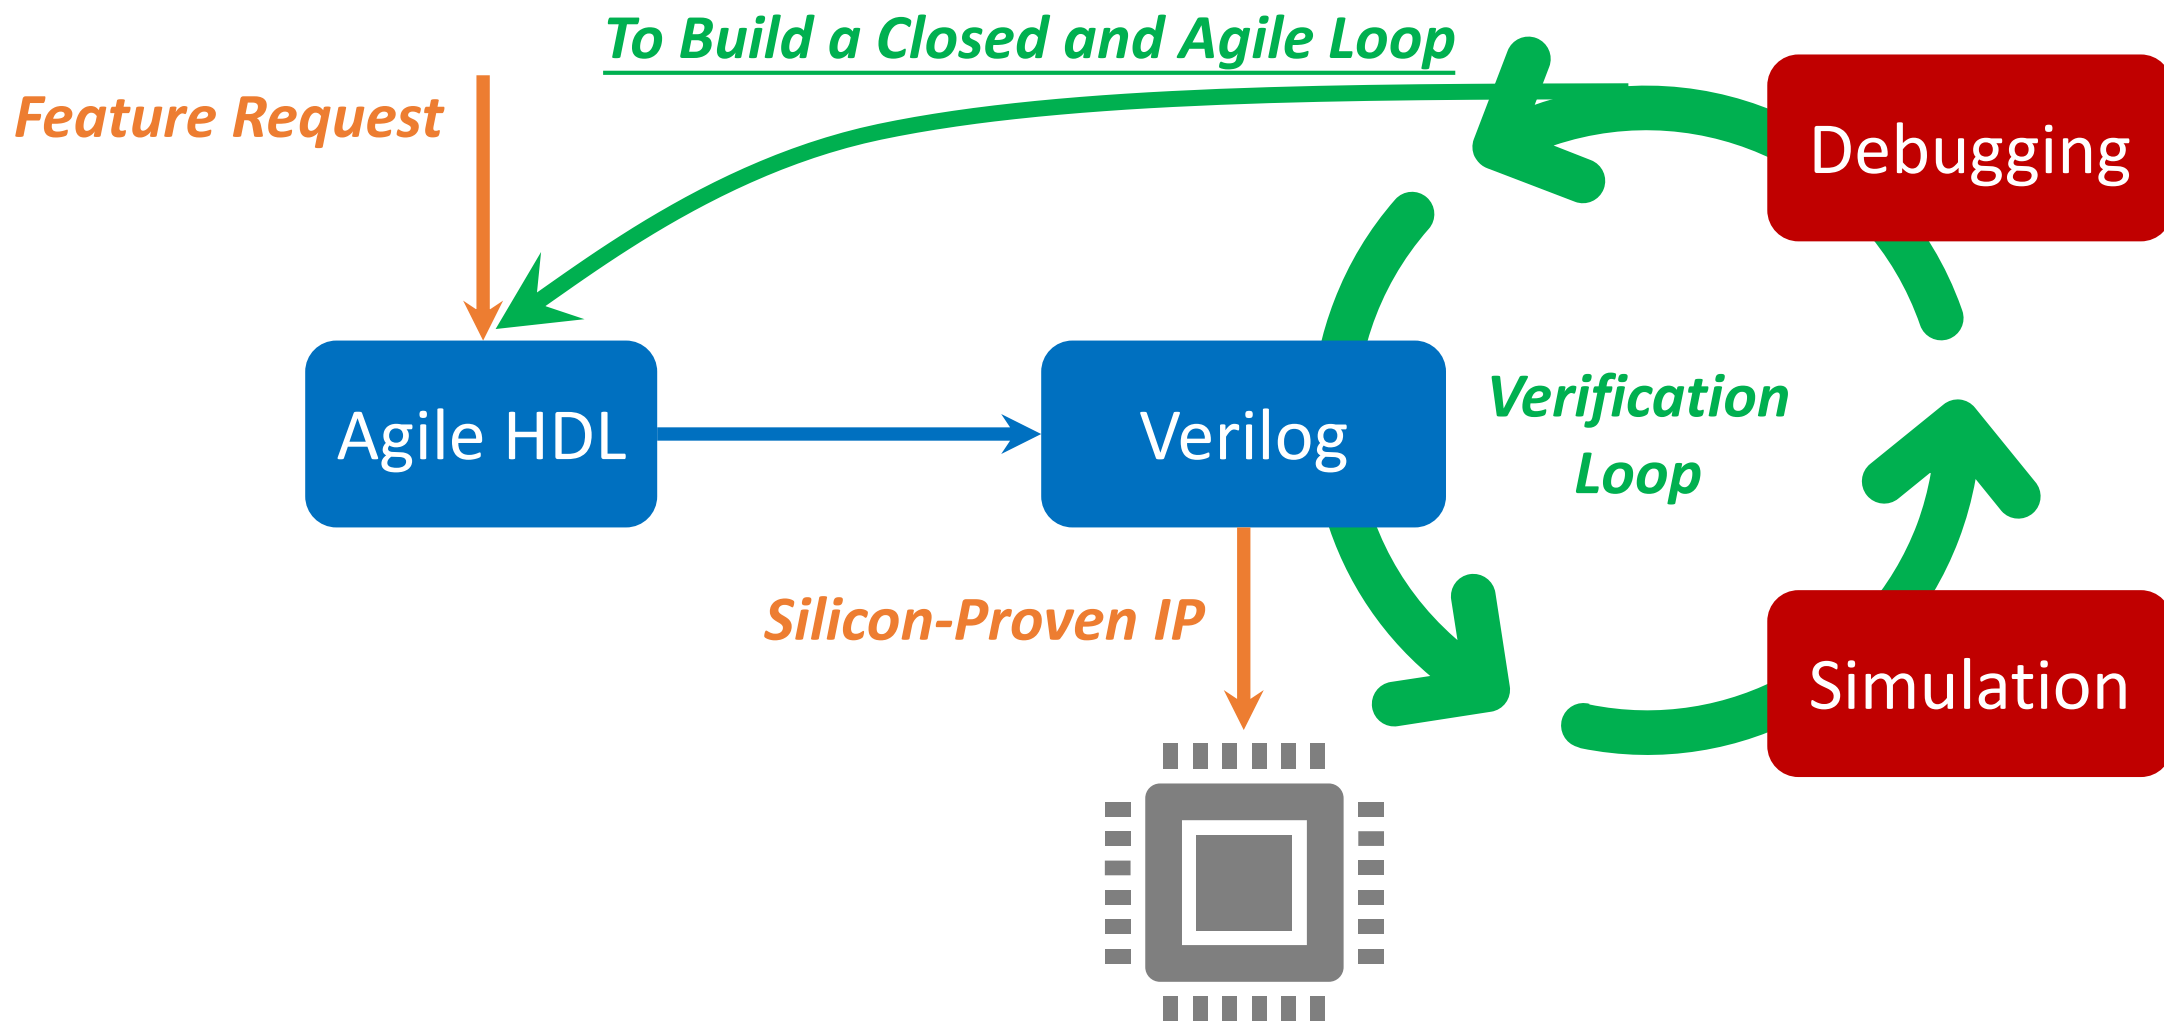


MinJie: Platform with Agile Development Flows and Tools



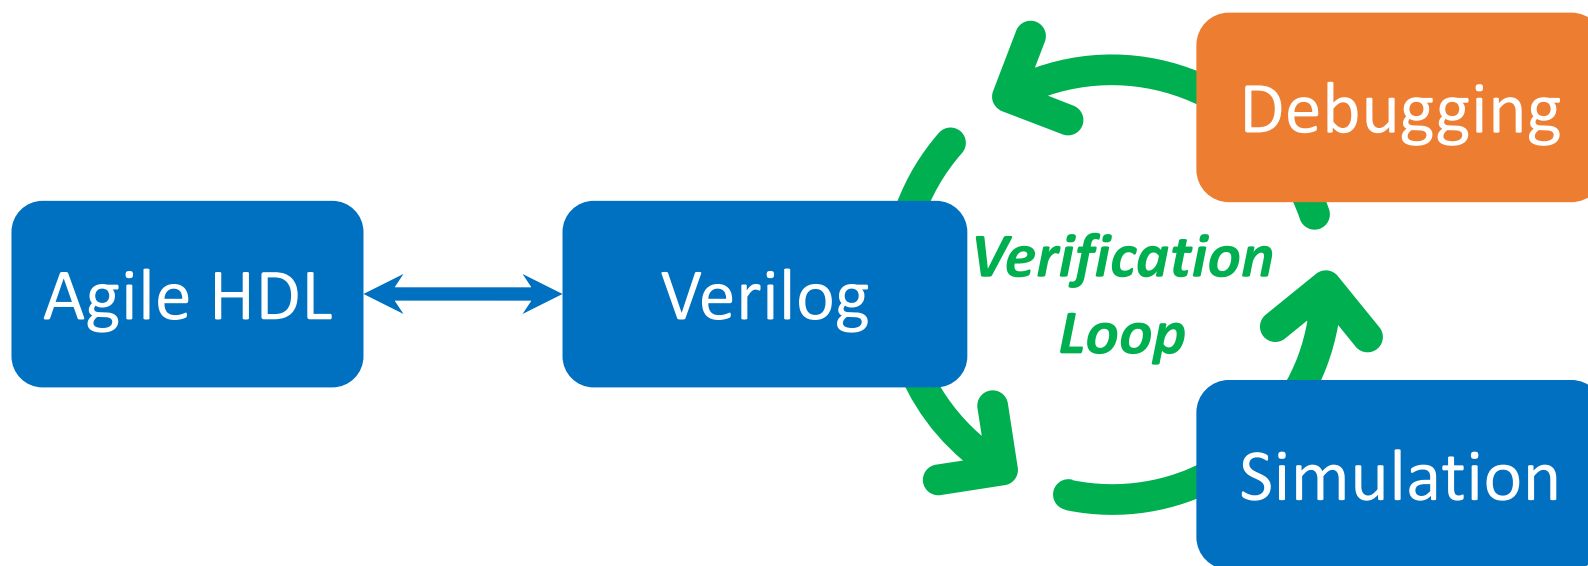


MinJie: Platform with Agile Development Flows and Tools



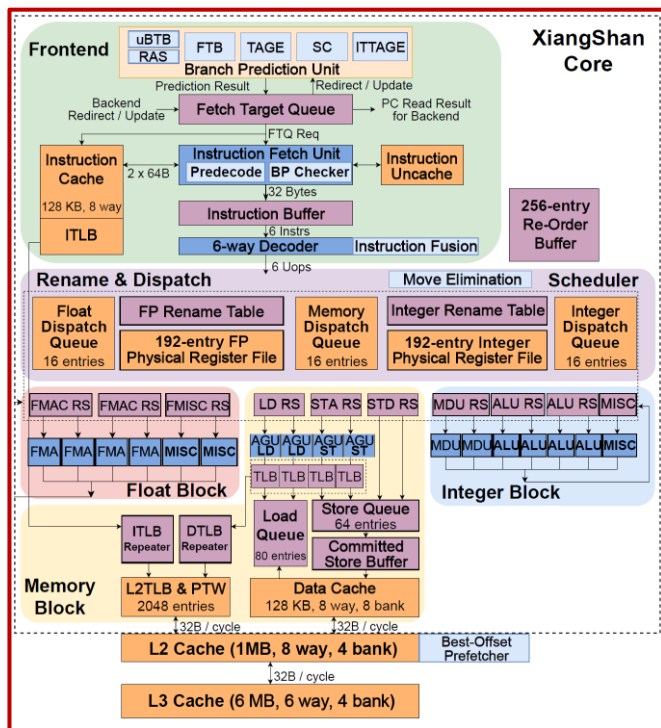
This Work

- *XiangShan*: High Performance RISC-V Processor
- *MinJie*: Platform with Agile Development Flows and Tools
 - *DRAV and DiffTest: Functional Verification*
 - *LightSSS: Debugging*



Processor Functional Verification

RISC-V Processor



The RISC-V Instruction Set Manual Volume I: Unprivileged ISA Document Version 20191213

Editors: Andrew Waterman¹, Krste Asanović^{1,2}
¹SiFive Inc.,
²CS Division, EECS Department, University of California, Berkeley
andrew@sifive.com, krste@berkeley.edu
December 13, 2019

RISC-V Specifications

The RISC-V Instruction Set Manual Volume II: Privileged Architecture Document Version 20211203

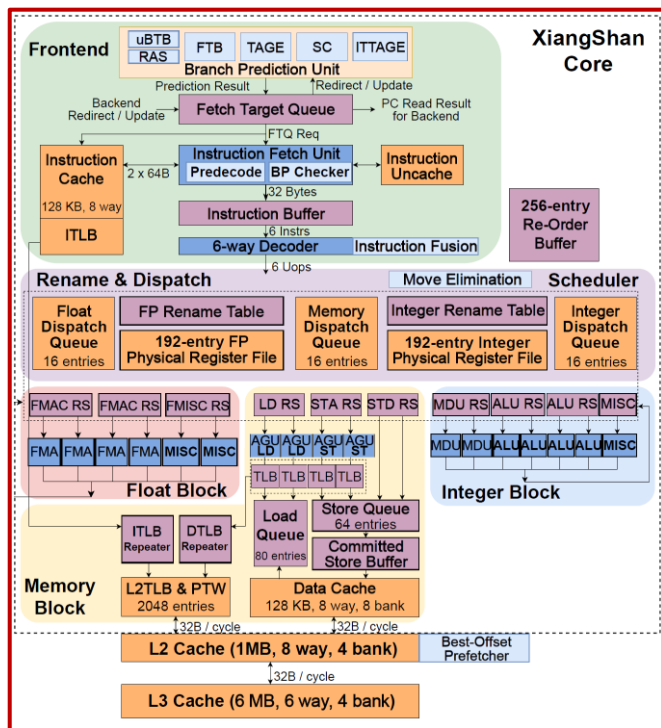
Editors: Andrew Waterman¹, Krste Asanović^{1,2}, John Hauser
¹SiFive Inc.,
²CS Division, EECS Department, University of California, Berkeley
andrew@sifive.com, krste@berkeley.edu, jh.riscv@jhauser.us
December 4, 2021

An ideal verification world:





Processor Functional Verification: Reality



**RISC-V
Processor**



**Yet Another
RISC-V
Processor**



**Non-Executable
RISC-V
Specifications**

SOTA: Let's co-simulate them and compare the results!

The RISC-V Instruction Set Manual
Volume I: Unprivileged ISA
Document Version 20191213

Editors: Andrew Waterman¹, Krste Asanović^{1,2}
¹SiFive Inc.,
²CS Division, EECS Department, University of California, Berkeley
andrew@sifive.com, krste@berkeley.edu
December 13, 2019

The RISC-V Instruction Set Manual
Volume II: Privileged Architecture
Document Version 20211203

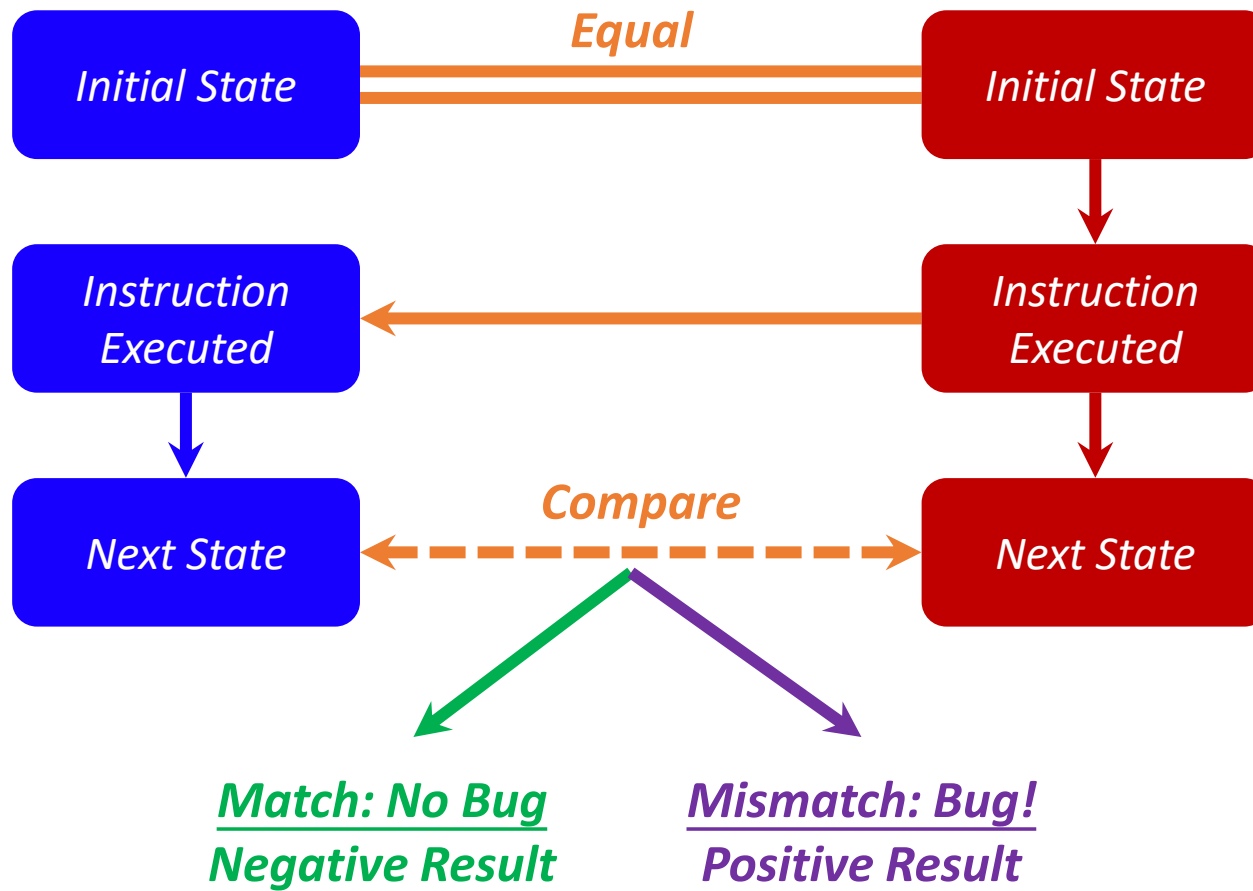
Editors: Andrew Waterman¹, Krste Asanović^{1,2}, John Hauser
¹SiFive Inc.,
²CS Division, EECS Department, University of California, Berkeley
andrew@sifive.com, krste@berkeley.edu, jh.riscv@jhauser.us
December 4, 2021



Processor Co-simulation to Find Bugs

RISC-V Model

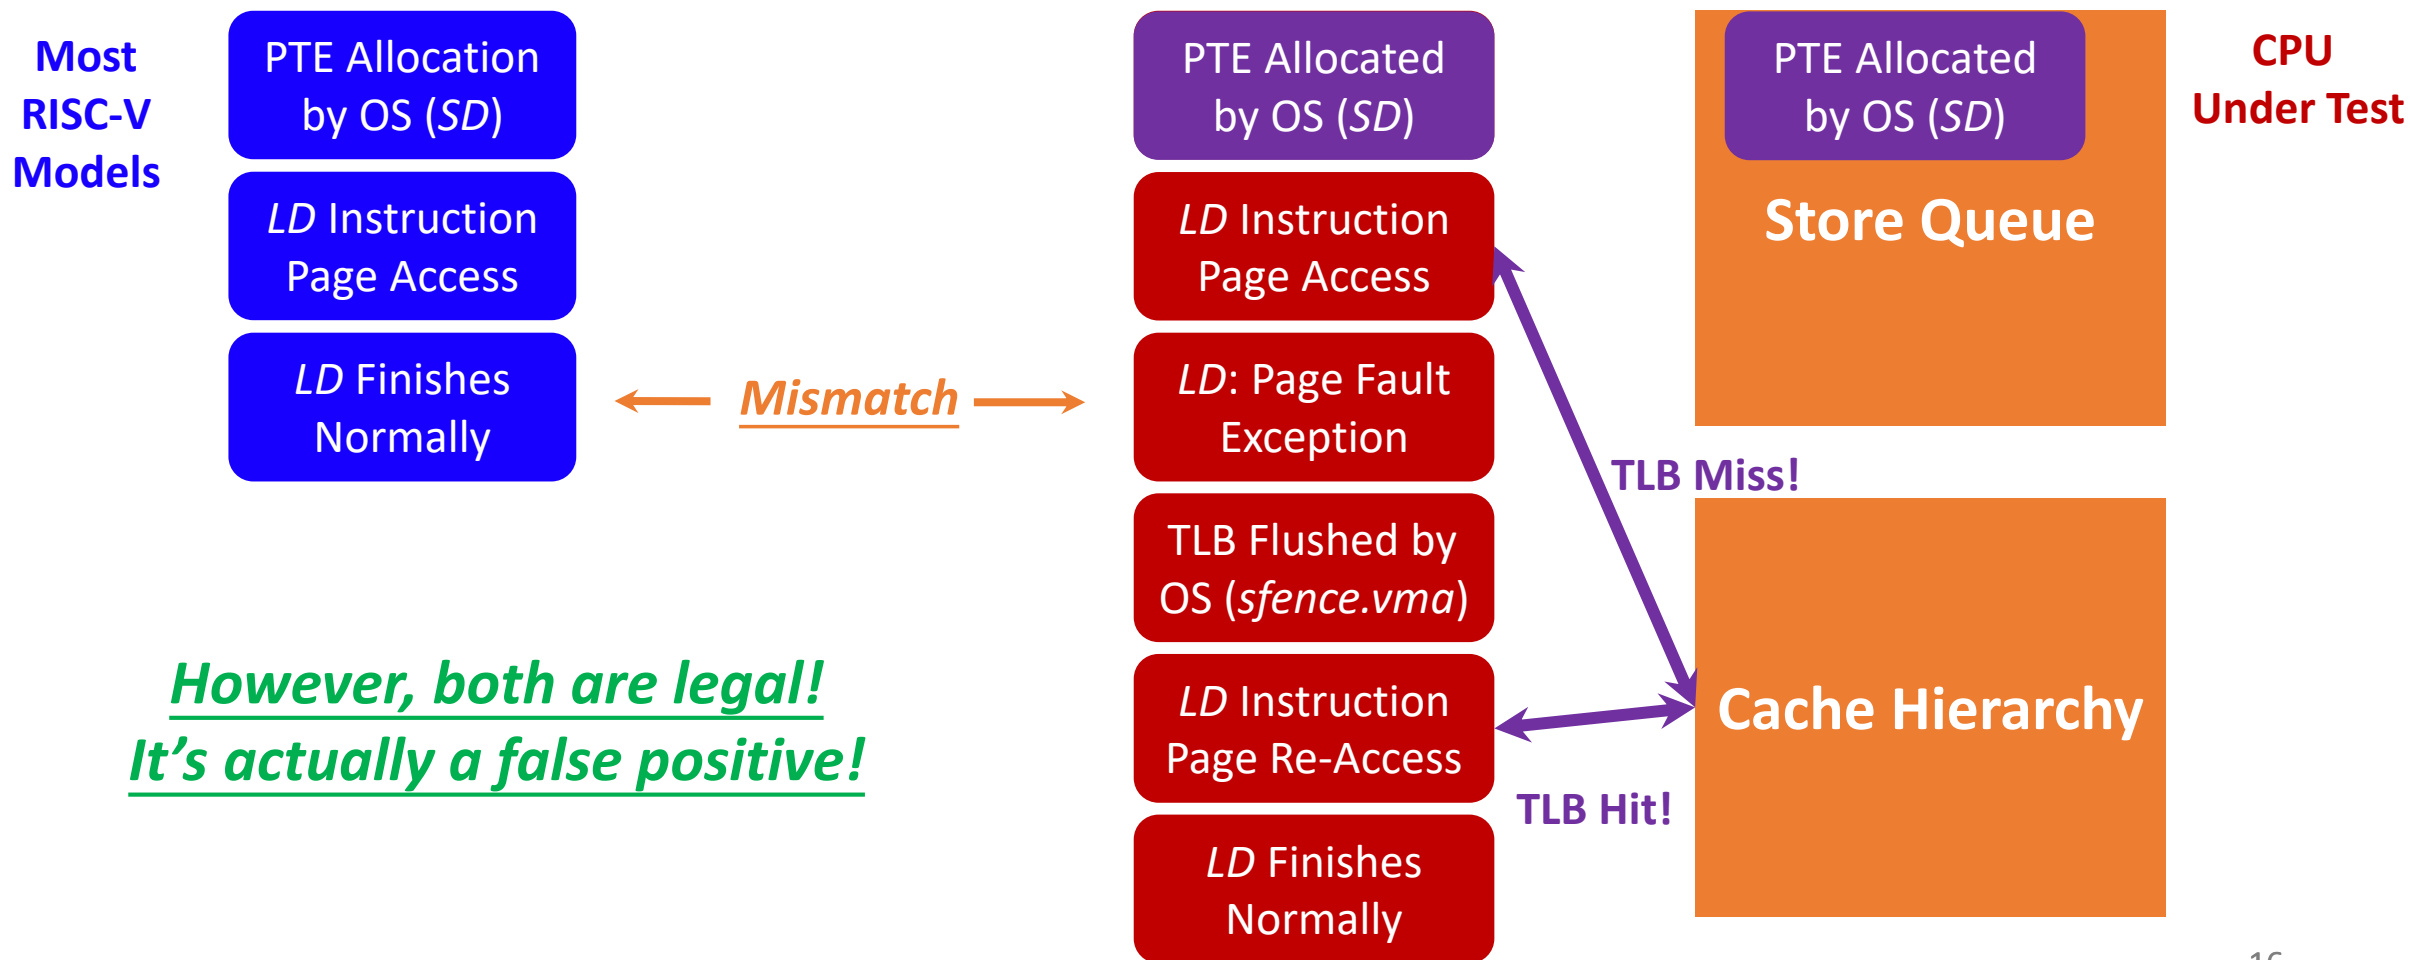
Processor Under Test





Challenge: False Positives

- **Example: Linux allocates valid PTEs and lazily executes memory-barrier instructions**
 - To avoid frequent TLB flushes for better performance






Observation: False Positives and Non-Determinism

- Designs are deterministic – as RTL/C++ code is fixed.



$$R_{P_i}: S_{P_i} \times E \rightarrow S_{P_i} \text{ for Processor } P_i$$

- Verification should be non-deterministic – as specifications allow diverse designs.

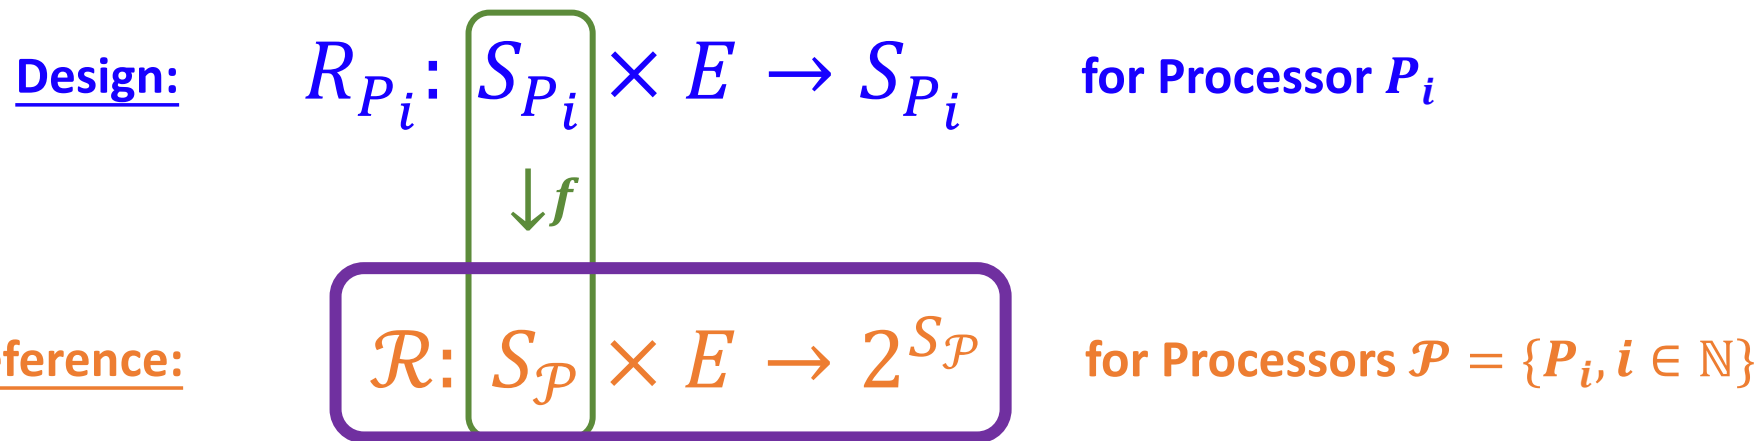


$$\mathcal{R}: S_{\mathcal{P}} \times E \rightarrow 2^{S_{\mathcal{P}}} \text{ for Processors } \mathcal{P} = \{P_i, i \in \mathbb{N}\}$$



This Work: *Diff-Rule based Agile Verification (DRAV)*

Probes: bridge hardware design and verification



Diff-Rules: abstract legal behaviors defined in specifications

DRAV: To Check Whether $f(R_{P_i}(s_{P_i}, e)) \in \mathcal{R}(f(s_{P_i}), e)$



Towards 1-to- N Correspondence Between REF and Designs

One REF for One Processor Design

Previously: $R_{P_i}: S_{P_i} \times E \rightarrow S_{P_i}$ for Processor P_i

Now: $\mathcal{R}: S_{\mathcal{P}} \times E \rightarrow 2^{S_{\mathcal{P}}}$ for Processors $\mathcal{P} = \{P_i, i \in \mathbb{N}\}$

One REF for N Processor Designs



This Work: *DiffTest* for RISC-V Processors

- Idea: acknowledge the non-deterministic nature of ISA – the ultimate golden model
- DiffTest: the state-of-the-art co-simulation framework for RISC-V processors
 - This Work: Identify and Specify Sources of Behavioral Non-Determinism in ISA using Diff-Rules
 - Dromajo[1]: to avoid non-deterministic sources such as the Debug Transport DTM (MMIO)
 - Imperas[2]: to extract asynchronous information from micro-architecture RTL pipeline

Categories	Sub-Categories	Examples	Addressed Before
Static	Impl. Dependent Registers	CSR, MMIO Devices	[1][2]
Dynamic	Impl. Dependent Registers	Timer, Counters	[2]
	Asynchronous Events	External/Timer Interrupts	[1][2]
	Speculative Execution	Instr./Load/Store Page Fault	<i>This Work Only</i>
	Memory Model	Memory Accesses in Multi-core	<i>This Work Only</i>
	Hardware Timing	LR/SC, Instruction Fusion, Caches	<i>This Work Only</i>

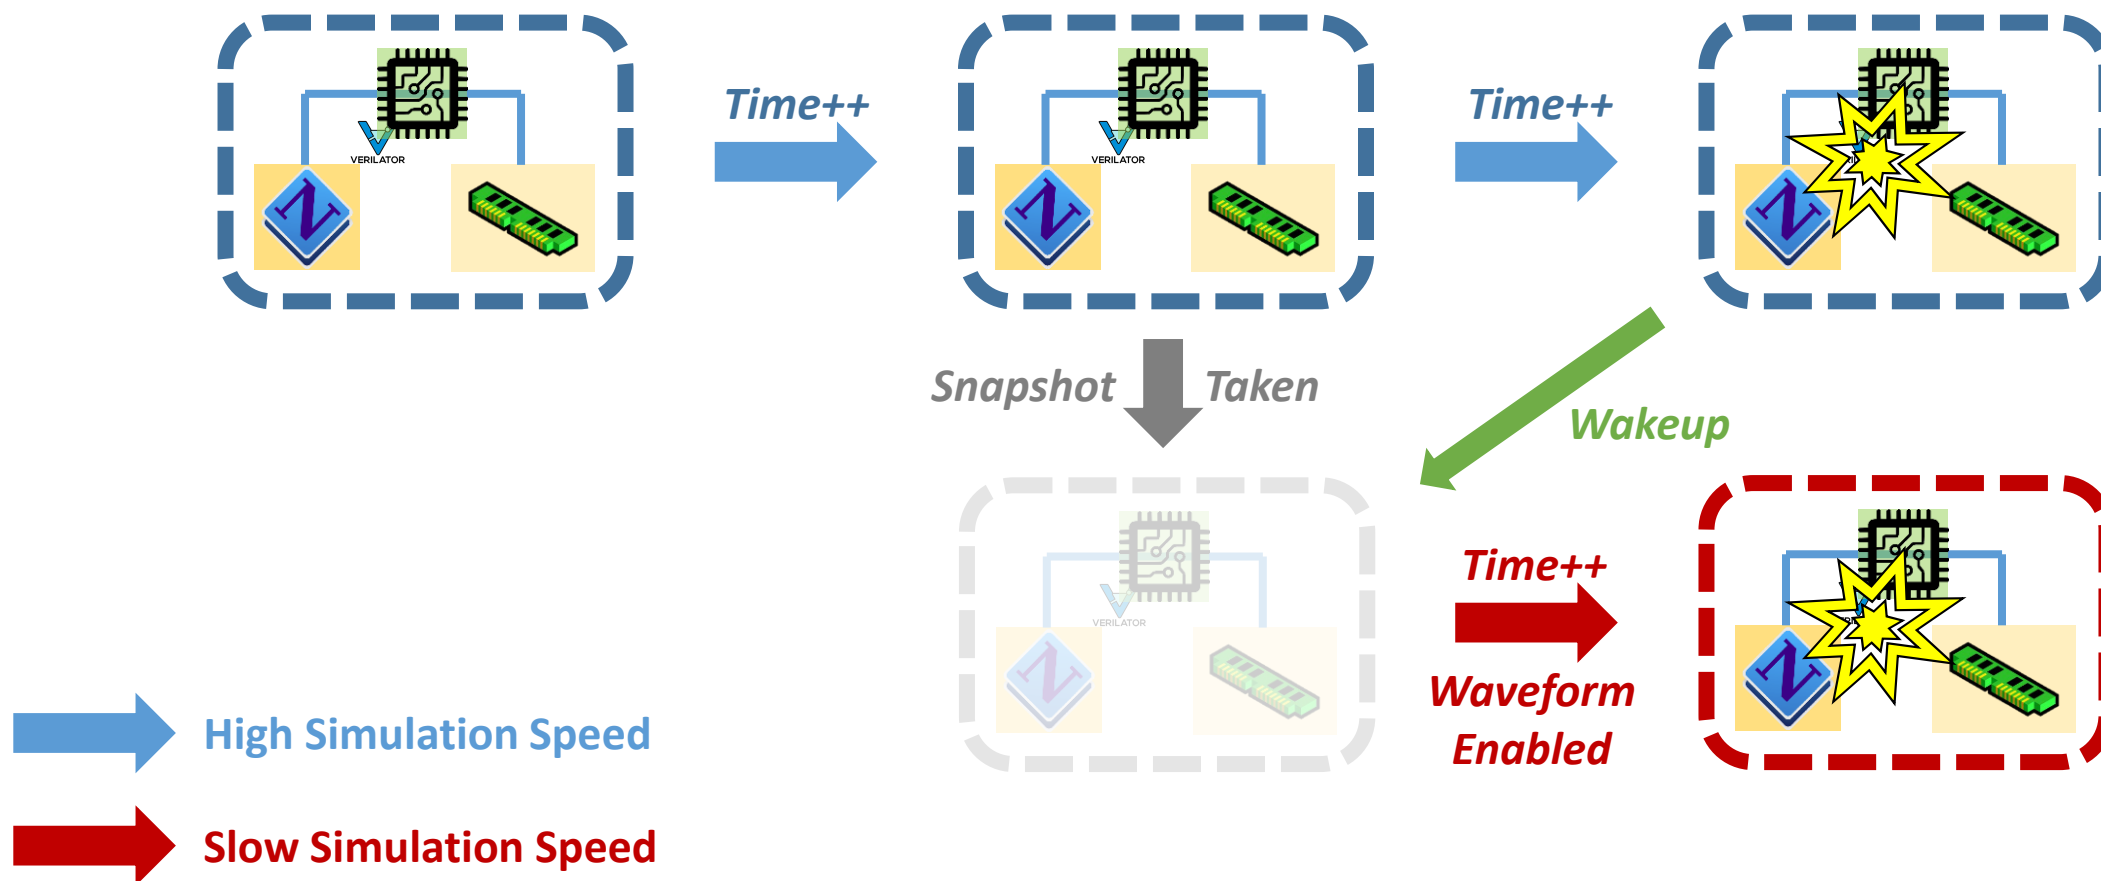
[1] Nursultan et al., Effective Processor Verification with Logic Fuzzer Enhanced Co-simulation. MICRO-54, 2021.

[2] Kevin McDermott. Brief Introduction to the 5 Levels of RISC-V Processor Verification. RISC-V Summit, 2021.



Accelerating Simulation Debugging with Snapshots


- Idea: to re-construct the last cycles of simulation after abort*





This Work: *Lightweight Simulation SnapShot (LightSSS)*

- Challenge: Minimize the Performance Overhead of Snapshots
- Key Insight: using fork from Linux to take snapshots of the simulation process
 - **Differential snapshot**: Copy On Write mechanism provided by the OS
 - **Portability**: transparent to the circuits, RTL simulators, external C/C++ models
 - **Efficiency**: in-memory snapshots without disk I/O

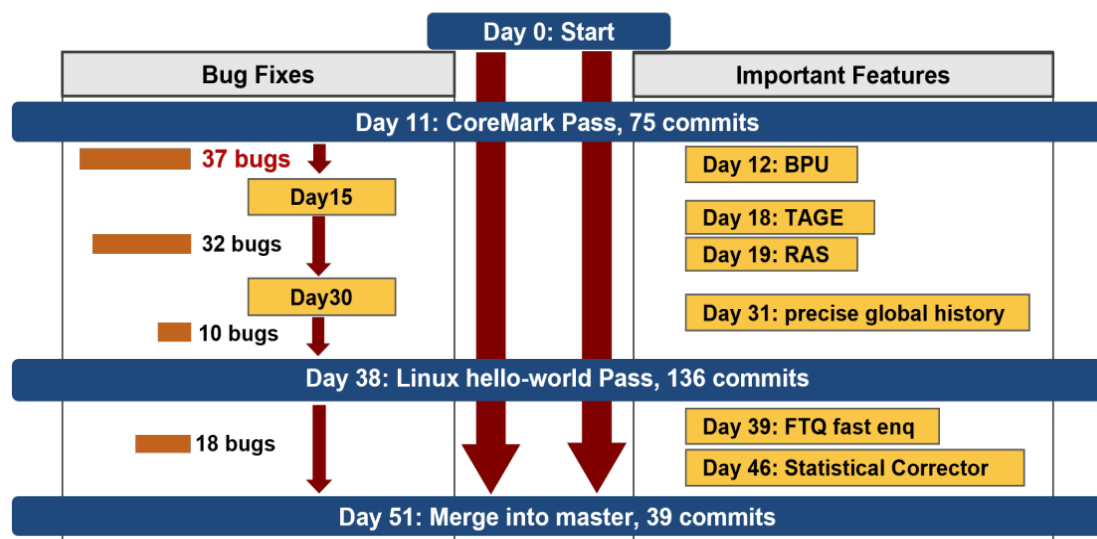


Minor performance overhead:
As snapshot interval size increases,
simulation speed remains stable



Practice of Agile Development on XiangShan

① Design Optimization





- 3 graduate students
- 11 days for a functionally correct prototype
- 37 bugs in 5 days
- 38 days to boot Linux with BPU
- 51 days for the overall frontend architecture

② Paper Reproduction

2018 51st Annual IEEE/ACM International Symposium on Microarchitecture

Performance Improvement by Prioritizing the Issue of the Instructions in Unconfident Branch Slices

One third-year PhD student, 200 minutes on XiangShan



More details in our paper

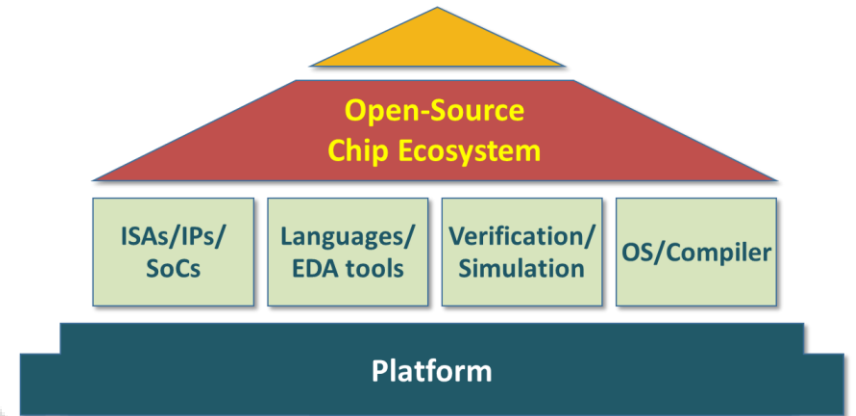

Conclusion

- **XiangShan: High Performance RISC-V Processor**
- **MinJie: Platform with Agile Development Flows and Tools**
 - DRAV and DiffTest: Functional Verification
 - LightSSS: Debugging
- **Both XiangShan and MinJie are Open-Source**
 - Follow us at [github/OpenXiangShan](https://github.com/OpenXiangShan) *with 3K stars* ❤️ *from the lovely community*
 - Artifacts Available, Functional, and Reproduced



Thanks!

Open-Source Chip Working Group



***Together for a Shared Future
To Lower the Barrier of Chip Development***



XiangShan Team



Open-Source Collaborators